

Shivansh Jain

India · +91 8959648831 · shivanshjain45@gmail.com · github.com/Shivansh-Jain · linkedin.com/in/shivansh-jain-in

SUMMARY

AI/ML Engineer with experience designing and scaling enterprise-grade AI platforms focused on Conversational AI, Generative AI, and Retrieval-Augmented Generation (RAG) systems. Led the architecture and development of scalable omni-channel conversational AI infrastructure supporting voice, chat, multilingual workflows, and real-time AI orchestration. Strong expertise in LLMs, agentic AI systems, distributed microservices, Kubernetes, Kafka, vector databases, and cloud-native deployments. Proven track record of building production-grade AI systems from proof-of-concept to enterprise-scale while leading engineering teams and driving technical strategy.

EXPERIENCE

Aivar Innovations *Senior AI/ML Engineer*

Bengaluru & Coimbatore, India | May 2025 – Present

- Promoted to Senior AI/ML Engineer within **6 months** for leading scalable enterprise conversational AI platform initiatives; awarded “**Thalaivar of the Year**” (Employee of the Year).
- Led architecture of **Convergent Accelerator** — transforming a telephony POC into an omni-channel conversational AI platform supporting SIP calling, WebRTC, WhatsApp/chat, multilingual conversations, and agent assist.
- Deployed microservices-based infrastructure on **Amazon EKS** with independently scalable services for orchestration, voice pipelines, analytics, and evaluation; integrated **KEDA** for dynamic autoscaling.
- Integrated **Kafka** for real-time event orchestration, conversational state management, CRM synchronization, and distributed workflow execution.
- Built a no-code flow-based conversational orchestration engine enabling enterprise users to configure workflows, branching logic, escalation paths, and AI behavior directly from the UI.
- Architected multilingual AI infrastructure with runtime switching between STT, TTS, and LLM providers including **ElevenLabs, Cartesia, Deepgram, Sarvam, AWS Bedrock, OpenAI**, and custom inference endpoints.
- Designed configurable **RAG pipelines** supporting Qdrant, OpenSearch, and PostgreSQL vector backends for domain-specific knowledge retrieval and hallucination reduction.
- Built advanced evaluation systems: persona-based testing, node coverage validation, AI-generated improvement suggestions, and real-time testing interfaces.
- Set up **OpenTelemetry + Grafana** for distributed tracing, production monitoring, and latency tracking; enabled both SaaS and fully self-hosted enterprise deployments.
- Led and mentored an AI engineering team of **10–12 engineers/interns**, driving engineering processes, deployment strategy, and architectural direction.
- Developed AI-powered freight extraction pipelines processing unstructured PDFs and Excel freight documents into structured JSON using OCR, table extraction, LLMs, and custom validation workflows.

Sustainability Economics *AI/ML Engineer*

Bengaluru, India | Aug 2024 – May 2025

- Created a RAG-based chatbot for the company’s main website, enabling contextual Q&A over proprietary knowledge bases.
- Developed an automated LLM-driven web data collection tool for domain-specific user queries.
- Built end-to-end RAG pipelines and vector database workflows for intelligent document retrieval and question answering.

Fiery (previously EFI) *Data Scientist Intern*

Bengaluru, India | Jan 2024 – Aug 2024

- Automated job property extraction from customer mail orders using **LLaMA 3.1** and Ollama, eliminating manual processing.
- Fine-tuned **Mistral-7B** with LoRA for improved language processing on internal domain-specific datasets.
- Built a RAG-based Q&A chatbot for company-specific knowledge retrieval using Mistral-7B.
- Developed vector database solutions for image and text similarity matching workflows.

Indian Institute of Technology Madras *Web Dev Intern*

Chennai, India | Jul 2023 – Dec 2023

- Developed a full-featured placement portal with PHP, JavaScript, and MySQL supporting admin, student, and recruiter workflows.
- Indian Institute of Technology Madras** *Academic Mentor — Python* Chennai, India | Aug 2022 – Sep 2022
- Mentored students in Python programming, guiding them from beginner to advanced proficiency.

EDUCATION

Indian Institute of Technology Madras *B.S. Data Science and Applications* Dec 2024
CGPA: 8.2 / 10

Vikram University, Ujjain *BBA, Business Administration* Mar 2022
79%

SKILLS

LLM & GenAI	LLaMA, Mistral, LoRA/QLoRA fine-tuning, Prompt Engineering, LangChain, Hugging Face, Transformers, Sentence Transformers, Ollama, OpenAI APIs
AI Systems	RAG Pipelines, Conversational AI, Agentic AI Workflows, Multi-Agent Systems, OCR & Document AI, Multilingual AI, Vector Search
ML / DL	PyTorch, Scikit-learn, NumPy, Pandas, SpaCy, Deep Learning, Sentiment Classification
MLOps & Infra	Amazon EKS, KEDA, Apache Kafka, Docker, Kubernetes, OpenTelemetry, Grafana, AWS Bedrock, ElevenLabs, Linux
Vector & Data	Qdrant, OpenSearch, PostgreSQL, MongoDB / AWS DocumentDB, Redis, MySQL, SQLite
Web & APIs	FastAPI, Flask, REST APIs, WebSockets, Vue.js, React.js, Node.js, Bootstrap, Jinja
Languages	Python (expert), JavaScript, PHP, SQL, Bash
Methodology	Conversational AI Architecture, Distributed Microservices, AI Workflow Orchestration, Scalable Inference Pipelines, AI Observability & Monitoring, Kubernetes-based AI Deployments

ACHIEVEMENTS & PROJECTS

- **Thalaivar of the Year — AIVAR** — Recognized as Employee of the Year for outstanding technical leadership and contribution toward scaling the company’s enterprise conversational AI ecosystem.
- **Sentiment Prediction on Movie Reviews** — Developed deep learning models for sentiment classification on movie review datasets.
- **Blog Lite 2.0** — Vue.js + Flask web app for blogging with likes, comments, and REST API backend (Flask-Restful, Flask-SQLAlchemy).
- **Blog Lite** — Multi-user blog sharing platform built on Flask with REST API and SQLAlchemy ORM.
- **Algoswarm Hackathon Website** — Designed and developed the official hackathon website using HTML, CSS, and JavaScript.